



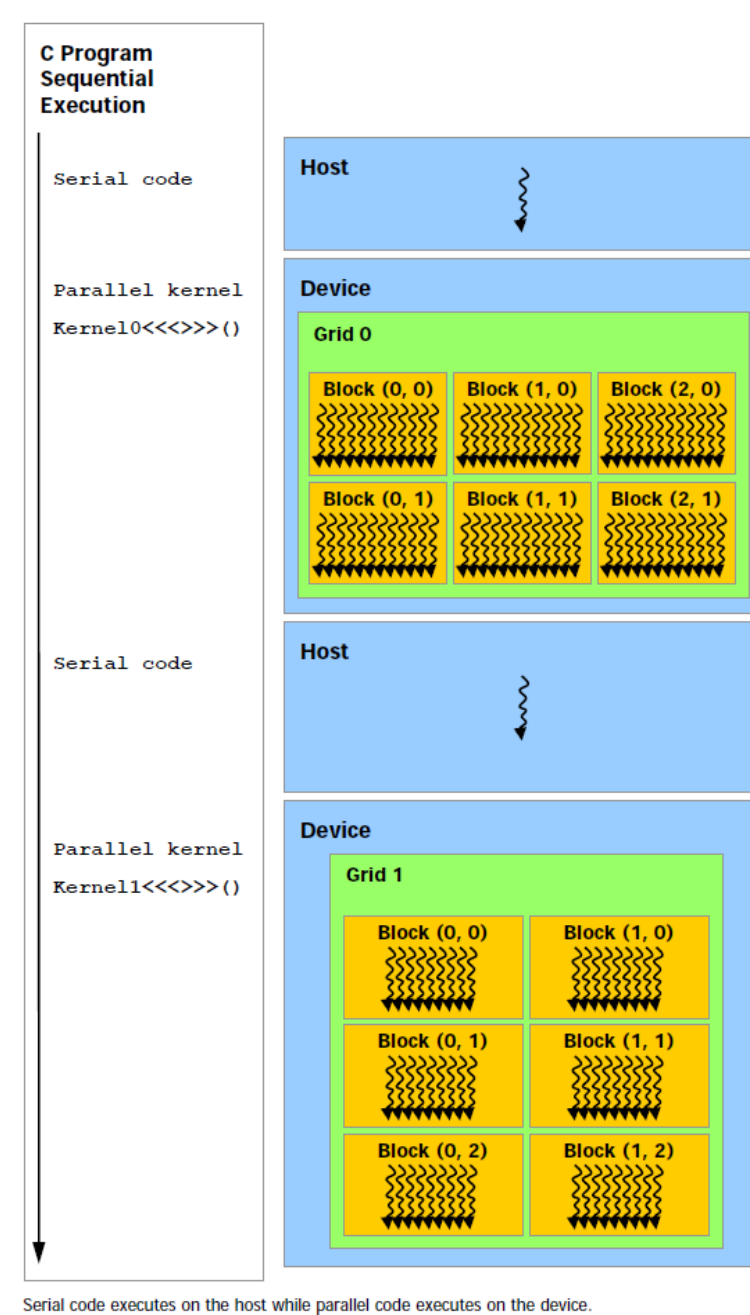
Goals & Motivations

- The iris pattern is considered to be amongst the most reliable for high confidence identification
 - Identification in enormous (national/global) databases
 - Requires ‘many-to-many’ comparisons; makes even the simplest distance metric computationally expensive
 - Fortunately: “embarrassingly parallel”
 - Distance matrix elements independent
- Standard solution: large CPU-based clusters
 - CPUs designed for general purpose, sequential tasks
 - Non-optimal w.r.t. power, cooling, footprint size & cost
- Better solution: acceleration with off-the-shelf GPUs
 - Market demand for realistic 3D games has evolved the GPU into a highly parallel, multithreaded, many-core processor of tremendous power



Compute Unified Device Architecture

- Prior to introduction of CUDA (API/architecture) in 2006 General-Purpose Computation on Graphics Hardware (GPGPU) was hard work:
 - Express problems in terms of graphics primitives
 - CUDA enables expression of programs in C, C++, Fortran and other high level languages
 - Heterogeneous: execute “kernels” on GPU
 - Section of device code that executes in parallel on GPU
 - Unique thread ID, program counter, registers and private local memory
 - Shared global device memory for communication
 - Hierarchy of memory types
 - SIMT model v. SIMD model
-



Results

	Template Size (bits)	GPU Exec Time HDs (ms)	GPU Exec Time Total (ms)	GPU Calc HD Rate (matches/s)	GPU Total Rate (matches/s)
57,109,864 Matches	(32 x 32) 1024	7576.78	7760.73	8,857,176.08	8,647,235.32
	(64 x 32) 2048	15936.50	16125.07	4,211,016.78	4,161,772.93
	(64 x 64) 4096	46968.48	47161.08	1,428,806.37	1,422,971.29
	(176 x 24) 4224	61575.40	17666.54	1,089,864.81	1,086,492.22
	(64 x 96) 6144	69043.42	69061.80	965,938.91	964,004.31
	(64 x 128) 8192	90406.05	90911.37	742,304.39	740,794.68

	Template Size (bits)	CPU ExecTime (ms)	CPU Matches/sec
Table Invariant	(32 x 32) 1024	127538.63	526.185
	(64 x 32) 2048	234918.5	268.560
	(64 x 64) 4096	454712.95	147.585
	(176 x 24) 4224	569861	117.764
	(64 x 96) 6144	675662.78	99.323
	(64 x 128) 8192	898720.51	74.672
SSE4.2 Invariant	(32 x 32) 1024	80453.74	834.130
	(64 x 32) 2048	138769.04	483.601
	(64 x 64) 4096	236847.72	283.342
	(176 x 24) 4224	317480.8	211.379
	(64 x 96) 6144	342103.57	196.165
	(64 x 128) 8192	443273.51	151.394

	Template Size (bits)	GPU Speedup (HD Only)	GPU Speedup (with memory)
Table Lookups	(32 x 32) 1024	16.8	16.4
	(64 x 32) 2048	14.7	14.6
	(64 x 64) 4096	9.7	9.6
	(176 x 24) 4224	9.3	9.2
	(64 x 96) 6144	9.7	9.7
	(64 x 128) 8192	9.9	9.9
SSE4.2 POPCOUNT	(32 x 32) 1024	10.6	10.4
	(64 x 32) 2048	8.7	8.6
	(64 x 64) 4096	5.0	5.0
	(176 x 24) 4224	5.2	5.1
	(64 x 96) 6144	4.9	4.9
	(64 x 128) 8192	4.9	4.9


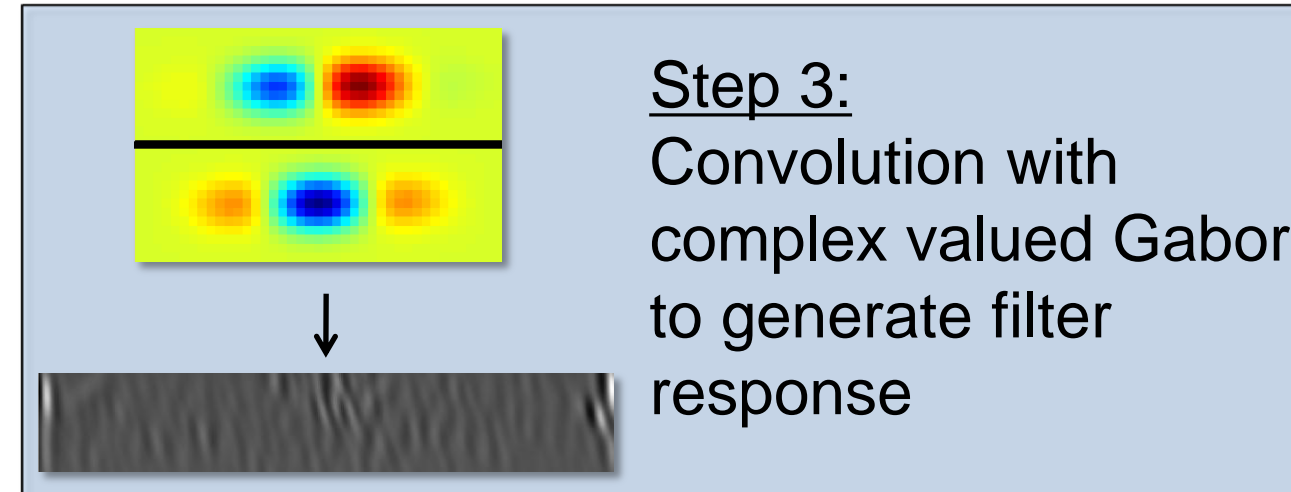


	Template Size (bits)	GPU Exec Time HDs (ms)	GPU Exec Time Total (ms)	GPU Calc HD Rate (matches/s)	GPU Total Rate (matches/s)
67100-864	(32 x 32) 1024	1515.70	1619.76	44,275.742	41,431,253
	(64 x 32) 2048	762.67	850.98	87,991.713	78,907,328
	(64 x 64) 4096	6126.10	6216.88	10,954.574	10,794,626
	(176 x 24) 4224	6684.06	6771.10	10,040.130	9,906.689
	(64 x 196) 6144	11072.07	11164.34	6,061.094	6,011,000
	(64 x 128) 8192	18272.66	19136.39	4,527.452	4,499,002

	Template Size (bits)	CPU Exec Time (ms)	CPU Matches/sec
Table Lookup	(32 x 32) 1024	9505.77	7,041,282
	(64 x 32) 2048	1823.42	3,684,565
	(64 x 64) 4096	4163.71	1,634,286
	(176 x 24) 4224	41594.70	1,613,399
	(64 x 96) 6144	60921.31	1,101,566
	(64 x 128) 8192	79505.10	844,083
POPCOUNT	(32 x 32) 1024	4282.87	15,669,151
	(64 x 32) 2048	8330.13	8,056,165
	(64 x 64) 4096	16033.12	4,185,641
	(176 x 24) 4224	16227.66	4,135,463
	(64 x 96) 6144	23438.37	2,863,206
	(64 x 128) 8192	37014.91	2,184,896


	Template Size (bits)	GPU Speedup (HD Only)	GPU Speedup (with memcpy)
Table Lookup	(32 x 32) 1024	12.5	11.2
	(64 x 32) 2048	12.0	11.4
	(64 x 64) 4096	6.7	6.6
	(176 x 24) 4224	6.2	6.2
	(64 x 96) 6144	5.5	5.5
	(64 x 128) 8192	5.4	5.3
SSE4.2 POPCOUNT	(32 x 32) 1024	5.6	5.0
	(64 x 32) 2048	5.5	5.2
	(64 x 64) 4096	2.6	2.6
	(176 x 24) 4224	2.4	2.4
	(64 x 96) 6144	2.1	2.1
	(64 x 128) 8192	2.1	2.1

Non-Rotation Invariant

How Iris Recognition Works



The most widely employed procedure for feature extraction, pioneered by John Daugman, uses the phase response of 2D Gabor wavelets.



$$g(x, y) = s(x, y) w_r(x, y)$$

$$s(x, y) = e^{j(2\pi(u_0 x + v_0 y) + P)} \quad w_r(x, y) = K e^{-\pi(a^2(x-x_0)_r + b^2(y-y_0)_r)}$$

$$(x-x_0)_r = (x-x_0)\cos\theta + (y-y_0)\sin\theta \quad (y-y_0)_r = -(x-x_0)\sin\theta + (y-y_0)\cos\theta$$

In addition to being a component of Daugman's algorithm, the template matching process (select the minimum fractional HD over a range of bitwise horizontal circular shifts) is a common final step of other iris recognition routines.

Implementation Overview

- Naïve implementation: exploit fine grain parallelism
 - Kernel below computes pairwise HD between a probe template and a gallery template determined by a unique two dimensional thread ID

```

__device__ float fractionalMaskedHD(unsigned int* x, unsigned int* y,
unsigned int* maskX, unsigned int* maskY, unsigned int len)
{
    unsigned int nB_c = 0;
    unsigned int nB_m = 0;
    unsigned int num;
    unsigned int* xEnd = x+len;
    unsigned int *u, *v, *mu, *mv;

    for(u = x, v=y, mu=maskX, mv=maskY; u!=xEnd;)
    {
        num = *mu++ + *mv++;
        nB_m += popc(num);
        num = (*u++ ^ *v++) & num;
        nB_c += __popc(num);
    }

    return __uint2float_rn(nB_c)/ __uint2float_rn(nB_m);
}

```

Listing 1: CUDA inline device code that computes single fractional Hamming distance using intrinsic function `__popc` to count set bits.

```

_global __void pdist2_smem_kernel(unsigned int* X, unsigned int* Y,
unsigned int* mX, unsigned int* mY, float D, unsigned int mx, unsigned int my,
unsigned int n)
{
    extern __shared__ unsigned int smem[];
    unsigned int* probeS = (unsigned int*)smem;
    unsigned int* maskS = (unsigned int*)smem[n];

    unsigned int y = IMUL(blockDim.y, blockIdx.y) + threadIdx.y;
    unsigned int tid = threadIdx.x;

    // Load the probe, which is constant within a given block into shared mem
    blockMemcpy(probeS, &X[blockIdx.x * n], tid, n, blockDim.y);
    blockMemcpy(maskS, &mX[blockIdx.x * n], tid, n, blockDim.y);

    __syncthreads();

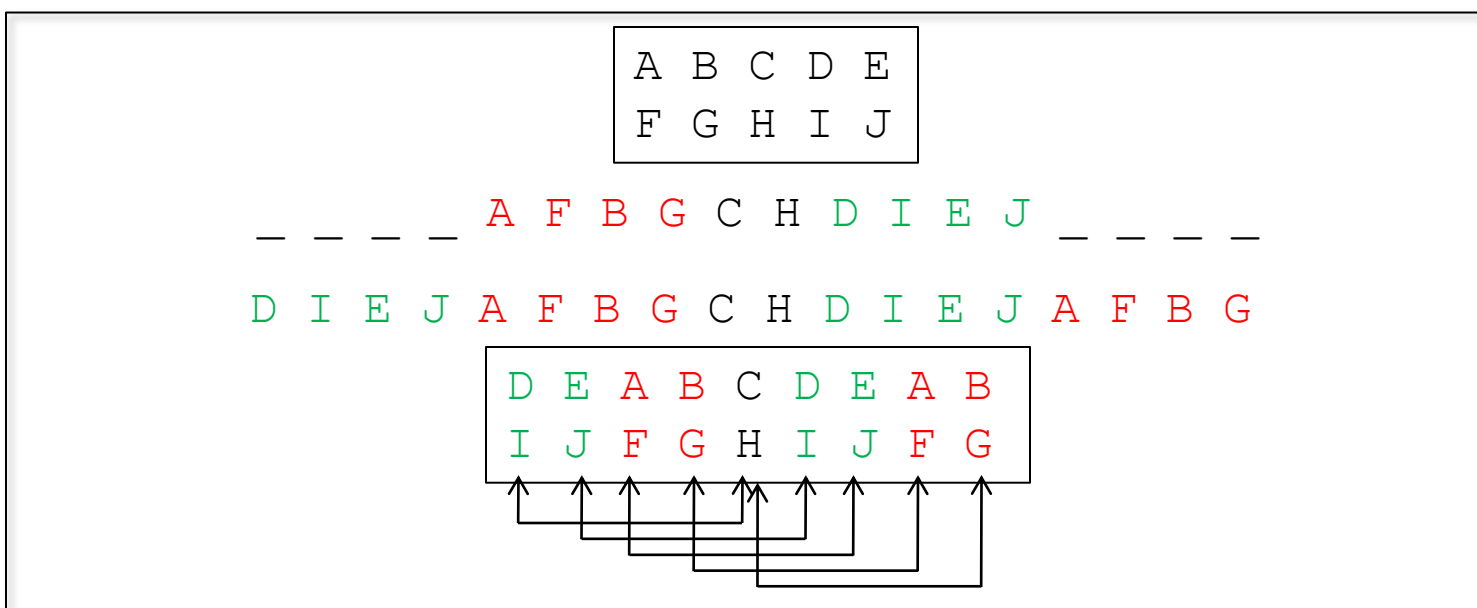
    if (y>=my)
        return;

    //Fractional HD
    D[Y[mx] + blockIdx.x] = fractionalMaskedHD(probeS, &Y[y*n], maskS, &mY[y*n], n);
}

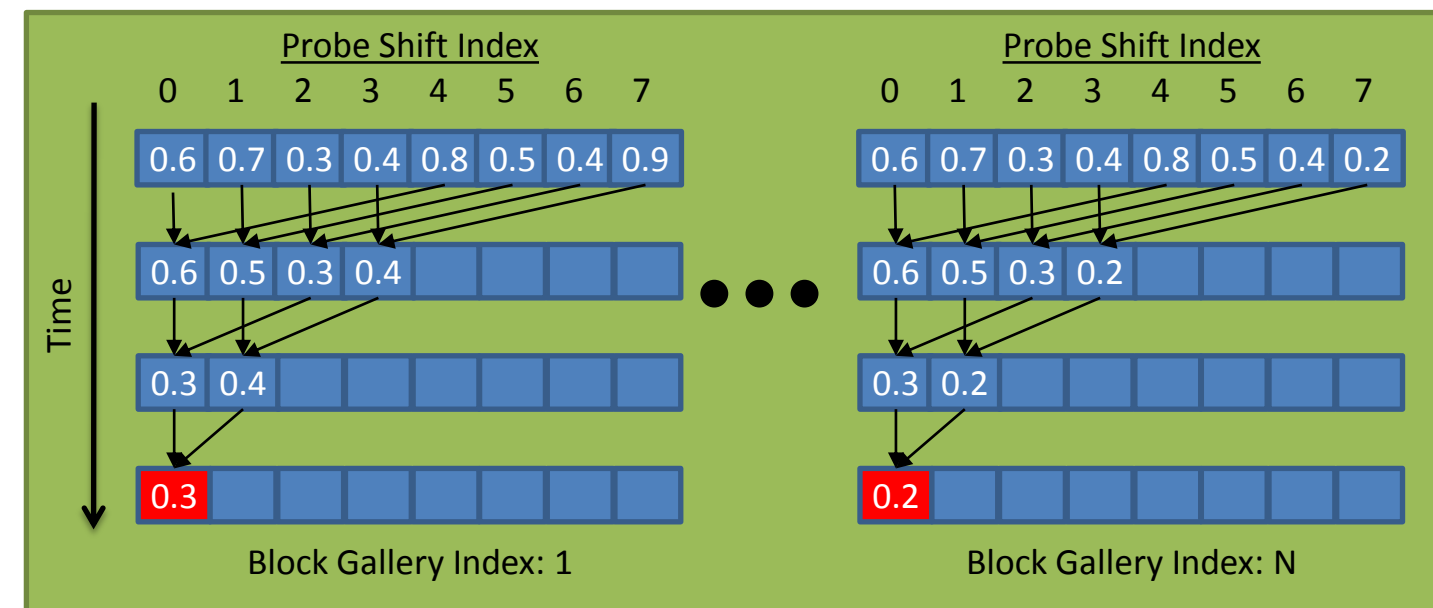
```

Listing 2: CUDA kernel code that utilizes shared memory, for nonshifted pairwise computation of fractional Hamming distance for set of probe and gallery templates.

- Optimizations: kernel is extremely memory bound
 - Exploit memory hierarchy \rightarrow shared memory and texture cache reduce global memory bandwidth usage/latency
- Rotation invariance:
 - Perform shifts in shared memory and reduce within thread block to find minimal HD for a given probe



Method using cooperative memcpy within a block's shared memory to execute all horizontal shifts from +2 to -2. Assumes column major matrix order with vertically packed bits



Simultaneous multi-row reduction performed in block shared memory used to find minimal fractional HDs over all possible shifts of probe.

Conclusions

- Achieved rates of 44 million iris template matches/s without rotation invariance. With tolerance to head tilt, 4.2 million matches/s (template size 2048 bits)
- Show a 14X speedup over optimized CPU implementation
- In contrast to other published work, our parallel implementation incorporates shifting for rotation invariance